

## Advantage

- Portable and flexible
- Use computer resources efficiently

## Disadvantage

- Significant overhead


## OS vs User Program

### Similarities

- Both softwares

### Difference

- OS runs in **kernel mode** - Access to all hardware resources
- User programs run in **User mode** - Limited access
- User programs use syscalls to communicate with OS for hardware processes
- User - Full virtual address space for code but cannot write outside of virtual address space
- Kernel - All code share same address space



- A OS executes machine instructions
- B Normal machine instructions executed
- C Calling OS using **syscall interface**
- D User programs call library code
- E System processes providing high level services

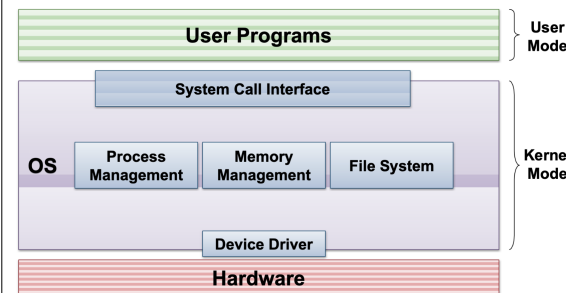
### Why OS dont occupy entire hardware layer

- Slow to have all operations pass through intermediary
- User programs can have direct interaction with hardware (eg. Arithmetic) during low risk operations

## OS structure


### Monolithic OS

- One big kernel program
- Well understood and has good performance
- Highly **coupled** - internal structure interconnected that unintentionally affect each other




## Microkernel

- Small clean
- Basic and essential facilities
- IPC communication OR run external programs outside OS
- Robust and more **modular** - Extendible and maintainable
- Better isolation btw kernel and services
- Lower performance
- Unix is monolithic kernel, Windows is hybrid



## Virtual Machines

- Software emulation of hardware
- Virtualization of underlying hardware to run additional operating systems concurrently

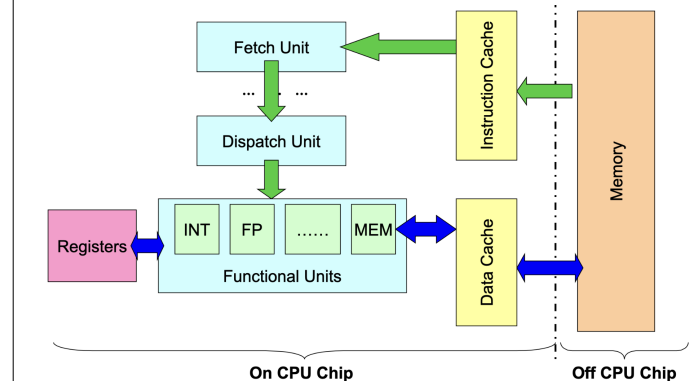


## 02. Process abstraction

### Motivation

- Allow concurrent usage of hardware
- Multiple programs sharing the same processors/IO


### Computer organisation



## 01. Introduction

**OS** - Program that acts as an intermediary between user and hardware

### Different architectures



**Difference** Separate vs common storage pathway for code and data

Why do we need OS?

- Should not have bugs
- Treat process as malicious

## Mainframe

Old analog "computers" using physical cards for programming

### Improvements

- Problem: Batch processing inefficient
  - Loading multiple jobs that runs while other jobs using I/O
  - Overlapping computation with I/O
- Solution: Multiprogramming
- Problem: Only one user
  - Multiple concurrent users using terminals
  - User job scheduling
  - Memory management
- Solution: Time sharing OS
  - **Hardware virtualization** - Each program executes as if it had all resources

## Motivation

1. Abstraction
  - Hide low level details and present common, high-level functionality to users
2. Resource allocation
  - Allow concurrent usage of resource and execute programs simultaneously
  - Arbitrate conflicting request fairly and efficiently
3. Control programs
  - Restrict resource allocation
  - Security, protection and error prevention (Defensive)
  - Ensure proper use of device

Memory

- Storage for instruction and data
- Managed by the OS
- Normally accessed via load/store instructions

Cache

- Fast and invisible to software
- Duplicate part of the memory for faster access
- Usually split into instruction and data cache

Fetch

- Load instructions from memory
- Location indicated by **Program Counter**

Functional units

- Carry out instruction execution
- Dedicated to specific instruction type

Registers

- Internal storage for fastest access speed

Information needed

- Memory context
  - Code
  - Data
- Hardware context
  - Register
  - PC value
  - Frame Pointer
- OS context
  - Process properties
  - Resources used
  - Files

Function calls

Suppose a function f() calls g()

- f is caller and g is callee

Steps of control flow

1. Setup parameters
2. Trf ctrl to callee
3. Setup local var

- Remember the initial variables especially if registers are to be reverted back when jumping back to caller
- Callee does not know the registers callers used, so prevent accidental overwrite of content to registers caller use

4. Store any results
5. Return ctrl to caller

Issues

Control Flow

- Need to jump to functional body when callee called
- Need to resume to next instruction in caller after done

Data storage

- Need to pass parameters to function
- Need to capture return result
- May have local variables

Additional

- May lead to overriding of data in caller by callee (interference)
- Calling g() multiple times may lead to insufficient space and overriding

Stack memory

Memory to store function invocation

**Stack Pointer** - Indicates the first free location in the stack region

**Frame Pointer** - Points to the frame and is used for traversing around the stack easily

■ On executing function call:

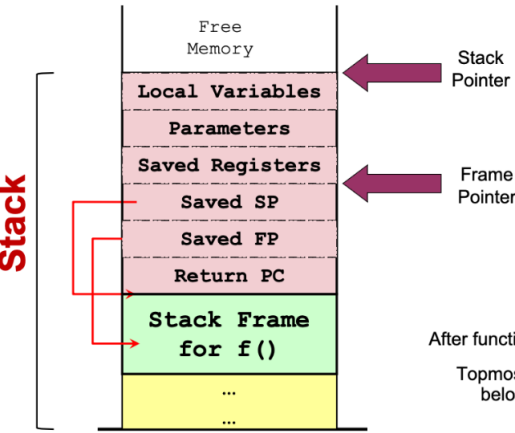
- **Caller:** Pass arguments with registers and/or stack ✓
- **Caller:** Save Return PC on stack
- **Transfer control from caller to callee**
- **Callee:** Save registers used by callee, Save old FP, SP
- **Callee:** Allocate space for local variables of callee on stack
- **Callee:** Adjust SP to point to new stack top

■ On returning from function call:

- **Callee:** Restore saved registers, FP, SP
- Transfer control from callee to caller using saved PC
- **Caller:** Continues execution in caller

Information needed for function invocation - Stack frame

- Return address of caller
- Arguments for the function
- Local variables
- Stack and frame pointer of caller
- GPR values (register spilling)



Callee stack frame will be on top of the caller

Dynamic memory (Heap)

Memory that the program/user specifies manually (eg. malloc, new)

Problems:

- Allocated only at runtime

- Size not known at program compilation time
- Cannot specify a region in data
- No definite dellocation timing
  - Must be freed explicitly by the program
  - Cannot place in stack region

Solution:

Add a region "Heap for dynamic allocation

Problems with heap memory:

- Generation of holes in between data due to variable deallocation timing

OS context

Process identification

Features:

- Distinguish processes from each other (Unique)
- Communicated to the hardware

Process state

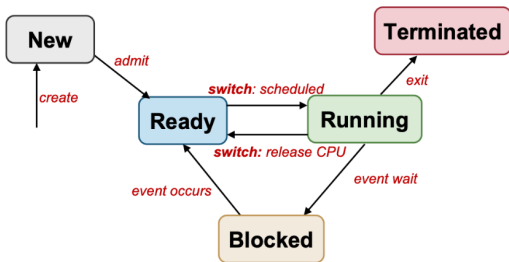
**New** New process that may still be under initialization (not schedulable)

Ready

**Running** Executed on CPU

**Blocked** Sleeping/waiting for event - cannot execute until event available

**Terminated** Finished execution and requires OS cleanup



<p><b>Process control block</b></p> <p>Table representing all processes containing entire execution context stored in OS memory</p> <ul style="list-style-type: none"><li>PC, FP, SP and other GPR (only updated when swapped out)</li><li>Memory region info (Text, data, heap and stack)</li><li>– It is a "point" to real memory space <b>NOT actual memory space used by process</b></li><li>PID and process state</li></ul> <p><b>Context switching</b></p> <ol style="list-style-type: none"><li>Interrupt or syscalls</li><li>Save state of old process into PCB</li><li>Reload state of new process from PCB</li></ol> <ul style="list-style-type: none"><li>Involves a change to kernel mode and back</li></ul> <p><b>Exceptions and interrupts</b></p> <p>Exceptions</p> <ul style="list-style-type: none"><li>Synchronous (due to program execution)</li><li>Machine level instructions arise errors</li><li>Exception handler executed automatically in software (can be handled by programmer)</li><li>Handled by OS which sends a signal to application to implement try-catch mechanism</li></ul> <p>Interrupts</p> <ul style="list-style-type: none"><li>Asynchronous (Can happen anytime)</li><li>External events that cause execution to fail (hardware related errors)</li><li>Program execution suspended and <b>interrupt handler</b> executed automatically</li></ul> <p><b>Instruction execution</b></p> <ol style="list-style-type: none"><li>Read byte from PC and decode instruction</li><li>Read 2 bytes to get the address/operands</li><li>Perform ALU operations</li><li>Store result into destination</li><li>Check if any interruptions</li></ol> <p>Interrupts can happen at anytime, and will remain pending until step 5 where it is handled</p>	<p><b>Interruption handling</b></p> <ol style="list-style-type: none"><li>Push PC and status register into hardware</li><li>Disable interrupts</li><li>Read <b>Interrupt Vector Table</b> - Table where the OS stores address of all interrupt handlers</li><li>Switch to kernel mode</li><li>Set PC to handler address and execute the instructions</li></ol> <ul style="list-style-type: none"><li>OS populates the IVT table with address of interrupt routines</li><li>Hardware reads IVT to locate the handler</li></ul> <p><b>System calls</b></p> <p><b>Application Program Interface</b> - Provides way of calling facilities/services in kernel</p> <p>Synchronous function only done in kernel mode providing system services</p> <p>Typically more expensive than library calls due to context switching</p> <p>OS dependent and cannot be avoided by any application</p> <p><b>Functions</b></p> <p><b>Process control</b> Direct processes like creating, terminating or synchronisation</p> <p><b>File/device manipulation</b></p> <p><b>Information maintenance</b> Get system data, time, etc</p> <p><b>Communication</b> Create, delete connections and send/receive messages</p> <p><b>Method (Programming language specific)</b></p> <ul style="list-style-type: none"><li>Library version with the same name and same arguments</li><li>User friendly library version</li><li>Using func <i>long syscall(long number);</i></li></ul>	<p><b>Mechanism</b></p> <ol style="list-style-type: none"><li>User invoke library call</li><li>Place call number in the designated location</li><li>Library call executes a special instruction (<b>TRAP/syscall</b>) to change user to kernel mode</li><li>(in kernel) syscall handler is determined (by a <b>dispatcher</b>)</li><li>syscall handler is executed</li><li>Syscall handler ends and control returned to the library call</li><li>Return to user mode and continue normal function mechanism</li></ol> <p><b>03. Process abstraction in Unix</b></p> <p>Process information</p> <p><b>Pid</b></p> <p><b>Proc state</b> Running, sleeping, stopped, zombie</p> <p><b>Parent pid</b></p> <p><b>Cumulative CPU time</b> For scheduling</p> <p><b>fork()</b></p> <p>Process creation</p> <p><b>Package</b> unistd.h and sys/types.h</p> <p><b>return</b> PID of newly created process(parent) and 0 (child process)</p> <p><b>Behaviour</b></p> <ul style="list-style-type: none"><li>Creates a child process</li><li>– <b>Copy</b> data of parent (Independent memory space)</li><li>– Unless explicitly modifying value at address, parent and child variables are distinct</li><li>– Sane code, same address space</li><li>– Differs by pid, ppid and fork() return value</li></ul>	<p><b>Implementation</b></p> <p>Clone the parent process</p> <ol style="list-style-type: none"><li>Create address space of child process</li><li>allocate new pid to child and pass to parent</li><li>Create kernel process data structures</li><li>Copy kernel environment of parent process</li><li>Initialize child process context (pid, ppid, cpu_time = 0)</li><li>Copy memory regions from parent</li><li>Acquire shared resources</li><li>Initialize hardware context for child process (copy parent registers)</li></ol> <p>Problem: Memcopy is very expensive operation</p> <p>Solution: Copy on write</p> <ul style="list-style-type: none"><li>Only duplicate a memory location when it is written to</li></ul> <p><b>exec(@param)</b></p> <p>Replace current executing process image</p> <ul style="list-style-type: none"><li>Code and data replaced</li><li>PID/PPID intact</li><li>If fail, program just continues to run the next instruction</li></ul> <p>Format</p> <p><b>param</b> char *path, char *arg0...</p> <ul style="list-style-type: none"><li>Note that last term <b>MUST</b> be <b>NULL</b> indicating end of argument list</li></ul> <p><b>header</b> unistd.h</p>
---	--	---	---

<b>exit(status)</b>
<b>return</b> Does not return anything
<b>param</b> int status - returned to the wait call
<ul style="list-style-type: none"> <li>Most system resource used by process are released on exit</li> <li>return from main() implicitly calls exit(0)</li> <li>Basic processes are not releasable <ul style="list-style-type: none"> <li>Pid and status</li> <li>Process accounting info</li> </ul> </li> </ul>
<b>wait(&amp;status)</b>
Parent child synchronisation
<b>param</b> &status - address to put return value
<b>header</b> sys/types.h and sys/wait.h
<b>return</b> pid of terminated process
<ul style="list-style-type: none"> <li>Call is blocking - suspend operation until at least one child terminates</li> <li>Cleans up remainder of child system resources (PID, status)</li> <li>waitpid - used for waiting for specific child process</li> </ul>
<b>Orphan and zombie process</b>
<b>Zombie</b> - All processes that has exited but parent did not call wait
<b>Orphan</b> - Child process whose parent has been terminated
<ul style="list-style-type: none"> <li>Parenthood will be propagated up to init which may use wait() to clean up</li> </ul>
<b>04. Inter Process Communication</b>
<b>Mechanism</b>
<ul style="list-style-type: none"> <li>Shared memory</li> <li>Message passing <ul style="list-style-type: none"> <li>Pipes</li> <li>Signal</li> </ul> </li> </ul>

<b>Shared memory</b>
Communication through read/write to shared memory
Advantages
<b>Efficient</b> Only require OS to setup shared region once
<b>Ease of use</b> Simple reads and write to memory <ul style="list-style-type: none"> <li>Implicit communication</li> <li>Can store any type of information</li> </ul>
Disadvantages
<b>Limited to single machines</b> Less efficient over different system
<b>Need Sync</b> Might have data races without synchronisation
<b>Race condition</b> - System behaviour is dependent on the context/interleaving of process → unpredictable outcome
<b>Steps</b>
<ol style="list-style-type: none"> <li>Create/locate a shared memory region M</li> <li>Attach M to process memory space</li> <li>Read from/write to M</li> <li>Detach M from memory space after use</li> <li>Destroy M <ul style="list-style-type: none"> <li>Only one process need to do this</li> <li>Can only destroy if M is not attached to any process</li> </ul> </li> </ol>
<b>Message Passing</b>
Explicit communication through exchange of message
<b>Naming</b> Have to identify the parties in the communication
<b>Synchronisation</b> Behaviour of sending/receiving operations
<ul style="list-style-type: none"> <li>Messages have to be stored in kernel memory space</li> <li>All sending/receiving operations have to be done through syscalls</li> </ul>

<b>Direct communication</b>
Sender/receiver explicitly name parties in communication
<ul style="list-style-type: none"> <li>One link per pair of communicating processes</li> <li>Need to know identity of other party</li> </ul>
<b>Indirect communication</b>
Message storage ( <b>Mailbox/Port</b> )
<ul style="list-style-type: none"> <li>Can be shared among a number of processes</li> </ul>
<b>Synchronisation behaviour</b>
Blocking
<ul style="list-style-type: none"> <li>send and receive is blocked until a message has received/sent (other party is ready)</li> </ul>
Non-blocking (asynchronous)
<ul style="list-style-type: none"> <li>execute immediately and sends information to somebody OR returns empty handed</li> </ul>
Typically, receive is synchronous and send is asynchronous
BUT send just buffers the message and only sends the message when the receiver is receiving (no loss)
Message buffers
<ul style="list-style-type: none"> <li>Under OS control</li> <li>Decouples sender and receiver - less sensitive to variation in execution</li> <li>Mailbox capacity declaration in advance</li> </ul>
<b>Rendezvous</b>
<ul style="list-style-type: none"> <li>Synchronous message passing</li> <li>Sender is blocked until receiver sends matching receive</li> <li>No buffering needed</li> </ul>
<b>Pros</b>
<ul style="list-style-type: none"> <li>Applicable beyond a single machine</li> <li><b>Portable</b> - Easily implemented on many platforms and processing environments</li> <li>Easier synchronisation <ul style="list-style-type: none"> <li>Implicit via send/receive behaviour</li> <li>Communication and synchronisation is done simultaneously</li> </ul> </li> </ul>

<b>Cons</b>
<b>Inefficient</b> Usually requires OS intervention every operation
<b>Difficult to use</b> Requires information packing into specific format
<b>Unix Pipes</b>
3 different communication channels to user
<ol style="list-style-type: none"> <li>stdin - standard in bounded to keyboard input</li> <li>stderr - used for error messages</li> <li>stdout - linked to screen</li> </ol>
” ” symbol for linking input/output channels of each process to each other
Create communication channel with 2 ends (one reading, one writing)
<ul style="list-style-type: none"> <li>Producer-consumer relationship</li> <li>Like anonymous file reading information FIFO</li> </ul>
<b>Synchronisation</b>
<ul style="list-style-type: none"> <li>Circular bounded byte buffer with implicit synchronization</li> <li>Writer wait when buffer is full</li> <li>Reader wait when buffer is empty</li> <li>Unidirectional (one write one read exclusively) vs bidirectional (any end for reading and writing)</li> </ul>
<b>pipe(@param)</b>
<b>param</b> array of file descriptors
<b>return</b> 0 for success, != for errors
<b>Unix signals</b>
Asynchronous notification about an event sent to process/threads
<ul style="list-style-type: none"> <li>Handle signal via default or user defined/supplied handlers(not all signals)</li> <li>Can only call async safe functions within handler (signal/wait cannot be called in POSIX)</li> <li>eg. Kill, stop, continue, error</li> </ul>

Standard signals			
First the signals described in the original POSIX.1-1990 standard.			
Signal	Value	Action	Comment
<b>SIGHUP</b>	1	Term	Hangup detected on controlling terminal or death of controlling process
<b>SIGINT</b>	2	Term	Interrupt from keyboard
<b>SIGQUIT</b>	3	Core	Quit from keyboard
<b>SIGILL</b>	4	Core	Illegal instruction
<b>SIGABRT</b>	6	Core	Abort signal from <a href="#">abort(3)</a>
<b>SIGFPE</b>	8	Core	Floating point exception
<b>SIGKILL</b>	9	Term	Kill signal
<b>SIGSEGV</b>	11	Core	Invalid memory reference
<b>SIGPIPE</b>	13	Term	Broken pipe: write to pipe with no readers
<b>SIGALRM</b>	14	Term	Timer signal from <a href="#">alarm(2)</a>
<b>SIGTERM</b>	15	Term	Termination signal
<b>SIGUSR1</b>	30,10,16	Term	User-defined signal 1
<b>SIGUSR2</b>	31,12,17	Term	User-defined signal 2
<b>SIGCHLD</b>	20,17,18	Ign	Child stopped or terminated
<b>SIGCONT</b>	19,18,25	Cont	Continue if stopped
<b>SIGSTOP</b>	17,19,23	Stop	Stop process
<b>SIGTSTP</b>	18,20,24	Stop	Stop typed at terminal
<b>SIGTTIN</b>	21,21,26	Stop	Terminal input for background process
<b>SIGTTOU</b>	22,22,27	Stop	Terminal output for background process

05.Process scheduling

Concurrency

- Multiple process progress at the same time
- Virtual parallelism (threading, context switching)
- Physical parallelism (multicore, multi CPU)

Scheduling

More ready processes vs limited CPU

- Each process diff CPU usage
- **Scheduling algorithms** - Deciding the process to run based on environment and process behaviour
- Processes have phases of io-bound and cpu intensive activities

Processing environment (sorted by increasing user interaction and response times):

1. Batch
2. Interactive
3. Real time

Objectives

**Fairness** Should get fair share of CPU time n prevent **starvation**  
- processes never have access to CPU

**Utilization** Hardware is used all the time

**non-preemptive** - Process gives up resources voluntarily **pre-emptive** - Process given a fixed quota to run and is suspended for other processes

Timeline

1. Scheduler is triggered (OS take over)
2. Context switching can happen and current process information stored and placed back in queue
3. Pick process to run using algorithm and setup its context

Batch

Non-preemptive used dominantly cos minimal user interaction needed

Objectives

**Turnaround time** Total time till finish running **including waiting** (finish time - arrival time)

**Throughput** Rate of task completion

**Makespan** Total time to complete **ALL** tasks

**CPU utilization** Percentage of **time** CPU is used

Algorithms

1. First Come First Serve (FCFS)
  - Scheduled using FIFO queue based on arrival time (bad turnaround time)
  - Guarantee no starvation - number of tasks in front of task is strictly decreasing
  - Reordering can reduce waiting time
  - **Convoy Effect** - Not all hardware resources are used concurrently because a process blocks it
2. Shortest Job First (SJF)
  - Choose tasks with smallest total CPU time
  - Estimate CPU time for tasks in advance using the previous CPU-bound phases
  - $Estimated_{n+1} = \alpha Actual_n + (1 - \alpha) Predicted_n$
  - Starvation is possible as long tasks end up continuously waiting
3. Shortest Remaining Time (SRT)

Interactive system

Objectives

**Response time** Time between request and response by system

**Predictability** Decreased variation in response time

Preemptive scheduling

- Periodic scheduler with time interrupts
- **Time quantum** - Execution duration given to a process
  - Constant(regardless whether the process has given up processing halfway) OR Variable(full time quantum given to a process)
  -

Algorithms

1. Round Robin (RR)
  - Store tasks in FIFO queue and pick tasks to run until
    - (a) Fixed time slice elapsed
    - (b) Task gives up CPU voluntarily
    - (c) Task blocks
  - Task placed at end of queue to wait for next turn
  - Response time guaranteed bounded by num task time quantum
  - Timer interrupt needed
  - Choice of time quantum duration is important
    - Longer quantum: Better utilization but greater wait time
    - Shorter quantum: Bigger overhead/ lower utilization but shorter wait time
  - Poor performance when there are many jobs all exceeding time quantum - high overhead from context switching
2. Priority scheduling
  - Assign tasks processes priorities based on impt and select them
  - Starvation for low priority tasks
    - Solution:
    - Lower priority every time the process have executed
    - Give processes a time quantum
  - **Priority Inversion** - Lower priority tasks locks the resources used by a higher priority task leading to deadlock OR allowing lower priority tasks to run before
    - Solution:



- Temporarily increase priority of tasks that lock resources until it unlocks
- Low priority task inherit the priority of high priority tasks which is restored once unlocked

### 3. Multi-level Feedback Queue (MLFQ)

- Runs processes with higher priorities but lower priority when the job fully utilise its time quantum
- Processes that voluntarily gives up/blocks before time quantum retains its priority
- New processes have highest priority
- Processes with same priority run in RR
- Long CPU processes are starved as its priority becomes very low and cannot complete
  - Solution: Occasionally reset tasks to full priority
- Processes can exploit and voluntarily give up to retain its priority
  - Solution: Peg priority based on total CPU usage instead

### 4. Lottery scheduling

- Give out tickets to processes for various system resources and randomly pick winners to be granted the resources
- **Responsive** - Newly created processes can immediately join the lottery
- Good control - Process can be assigned quantity of tickets and resources which can be shared among its children/threads
- Simple implementation

## 06.Synchronisation primitives

### Race condition

- Execution of sequential processes should be **deterministic** - Repeated execution returns the same result... however,
- Process sharing modifiable resources AND execute concurrently by interleaving may cause synchronization problems - non deterministic outcomes
- Order determines the execution outcome

### Critical section

Region where unsynchronised access could lead to incorrectly interleaving scenarios

Only **ONE** process should be in the critical section at one time

### Properties

**Mutual exclusion/Mutex** All other processes should be blocked from entering critical section if occupied

**Progress** If no process occupying section, waiting process should be granted access

**Bounded wait** Upper bound on number of times other processes can enter critical section before a waiting process

**Independence** Process not in critical section should not block other process

### Incorrect synchronization

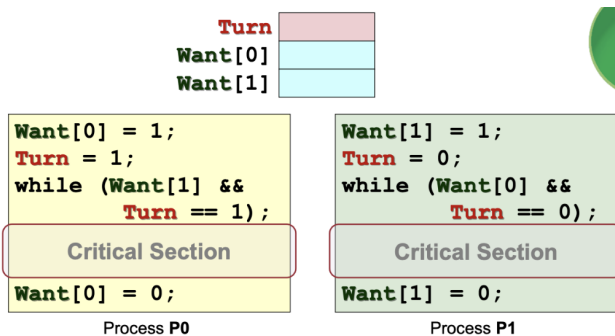
**Deadlock** All processes blocked

**Livelock** Processes keep changing states to avoid deadlock resulting in no progress (Deadlock avoidance mechanism)

### Starvation

### Implementations

- High level language implementation (Peterson Algorithm)
  - Busy waiting - Repeatedly tests while-loop condition instead of going into blocked state (Wastes CPU power)
  - Low level - Error prone and complicated
  - Not general - cannot extend to allowing multiple access instead of Mutex




- Assembly implementation

- Atomic instruction to load value and replace value in memory location (**TestAndSet**)
- **Atomicity** - Prevents race condition of process viewing and upon modification, the value changes
- No bounded wait unless fair scheduling algorithm employed
- Employs busy waiting still

```

void EnterCS( int* Lock )
{
    while( TestAndSet( Lock ) == 1);
}
  
```

- returns the current content of M  
- sets content of M to 1



```

void ExitCS( int* Lock )
{
    *Lock = 0;
}
  
```

- High level synchronization mechanism (Semaphore)

- Generalised mechanism to block a number of processes (sleeping),  $S$ , and unblock sleeping processes
- **wait()** - Blocks if  $S \leq 0$ , decrement  $S$
- **signal()** - Increment  $S$ , wake up sleeping processes, operation **never** blocks
- **Invariant** -  $S_{current} = S_{initial} + N_{signal} - N_{wait}$

### Conditional variables

- Tasks wait for certain events to happen
- Event completion/begin is broadcasted

### POSIX semaphore

header semaphore.h

wait pthread\_mutex\_unlock

signal pthread\_mutex\_lock

### Midterm

- **Tail-call optimization** - Stack frame can be replaced in iterative recursion since all information is retained in the 'return' function call at every iteration

## 07.Threads

### Motivation

- Processes are expensive
  - Process creation via fork() duplicates memory space and process context
  - Requires context switching and saving/restoring process information repeatedly

- Hard to communicate with each other
  - Independent memory space (IPC communication only)
- Need easy way to execute some instructions simultaneously (more threads of control to same process)

**Features**

- **Multithreaded process** - Single process with multiple threads
- Threads share the same **Memory context**(Text, data, heap) and **OS context**(PID)
- Contains unique information
  - Identification (thread id)
  - Registers (GPR and special)
  - Stack

**Benefits**

**Economy** Less resources to manage

**Resource sharing**

**Responsiveness**

**Scalability** Can take advantage of multiple CPUs

**Problems**

- Syscall concurrency
  - Parallel execution of threads may result in parallel syscalls
- Process behaviour
  - Unable to impact process operations

**Thread implementation**

1. User thread
  - Implemented as a user library - More flexible and configurable
  - Runs on any OS
  - Kernel unaware of threads in process (scheduling at process level)
  - If process is blocked, all the threads in the process cannot run
2. Kernel thread
  - Implemented in the OS through syscalls (slower and more resource intensive)

- Scheduling among threads instead of via process
- Use threads in kernel operations
- Less flexible
  - Used by all multithreaded programs so kernel must cater to all
  - Balance btw many and less features

3. Hybrid thread

- OS schedules on kernel threads
- User threads can be bound to kernel threads
- If the user thread is blocked, the kernel thread that it is assigned to is blocked and the OS will execute another kernel thread
- Lead to greater flexibility

**Posix Threads (pthreads)**

**pthread\_create**

**param** tid, attributes, &function, function\_params

**return** 0 success, != 0 fail

**pthread\_exit**

Automatically called after function finishes

**param** exit\_value

**behaviour** return value of function will be the "exit value"

**pthread\_join**

**param** tid, **&status** - Exit value returned by target pthread

**08.Synchronization classics**