


## 01. Introduction

**OS** - Program that acts as an intermediary between user and hardware

### Different architectures



**Difference** Separate vs common storage pathway for code and data

Why do we need OS?

### Mainframe

Old analog "computers" using physical cards for programming

### Improvements

- Problem: Batch processing inefficient
- Solution: Multiprogramming
  - Loading multiple jobs that runs while other jobs using I/O
  - Overlapping computation with I/O
- Problem: Only one user
- Solution: Time sharing OS
  - Multiple concurrent users using terminals
  - User job scheduling
  - Memory management
  - **Hardware virtualization** - Each program executes as if it had all resources

### Motivation

1. Abstraction
  - Hide low level details and present common, high-level functionality to users
2. Resource allocation
  - Allow concurrent usage of resource and execute programs simultaneously
  - Arbitrate conflicting request fairly and efficiently
3. Control programs
  - Restrict resource allocation
  - Security, protection and error prevention
  - Ensure proper use of device

### Advantage

- Portable and flexible
- Use computer resources efficiently

### Disadvantage

- Significant overhead


### OS vs User Program

Similarities

- Both softwares

Difference

- OS runs in **kernel mode** - Access to all hardware resources
- User programs run in **User mode** - Limited access
- User programs use syscalls to communicate with OS for hardware processes
- User - Full virtual address space for code but cannot write outside of virtual address space
- Kernel - All code share same address space



- A OS executes machine instructions
- B Normal machine instructions executed
- C Calling OS using **syscall interface**
- D User programs call library code
- E System processes providing high level services

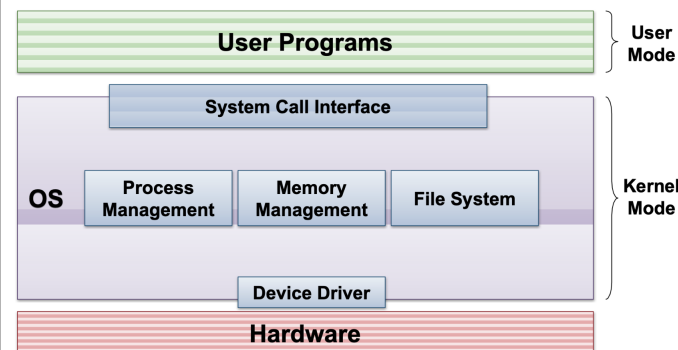
Why OS dont occupy entire hardware layer

- Slow to have all operations pass through intermediary
- User programs can have direct interaction with hardware (eg. Arithmetic) during low risk operations

### OS structure

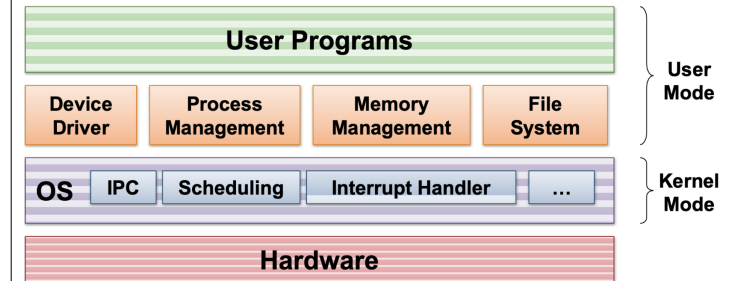
#### Monolithic OS

- One big kernel program
- Well understood and has good performance
- Highly **coupled** - internal structure interconnected that unintentionally affect each other




### Microkernel

- Small clean
- Basic and essential facilities
- IPC communication OR run external programs outside OS
- Robust and more **modular** - Extendible and maintainable
- Better isolation btw kernel and services
- Lower performance



### Virtual Machines

- Software emulation of hardware
- Virtualization of underlying hardware to run additional operating systems concurrently

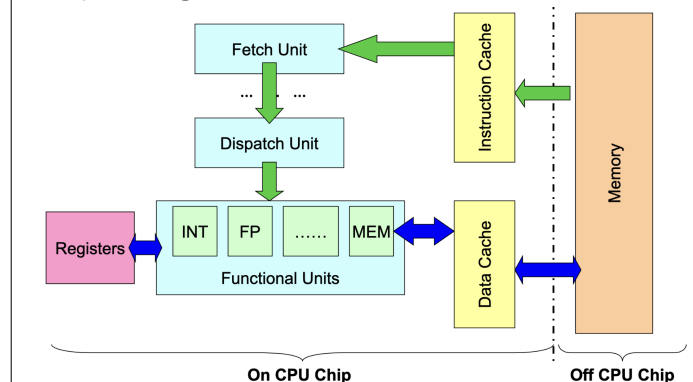


## 02. Process abstraction

### Motivation

- Allow concurrent usage of hardware
- Multiple programs sharing the same processors/IO

### Computer organisation



Memory

- Storage for instruction and data
- Managed by the OS
- Normally accessed via load/store instructions

Cache

- Fast and invisible to software
- Duplicate part of the memory for faster access
- Usually split into instruction and data cache

Fetch

- Load instructions from memory
- Location indicated by **Program Counter**

Functional units

- Carry out instruction execution
- Dedicated to specific instruction type

Registers

- Internal storage for fastest access speed

Information needed

- Memory context
  - Code
  - Data
- Hardware context
  - Register
  - PC value
  - Frame Pointer
- OS context
  - Process properties
  - Resources used
  - Files

Function calls

Suppose a function f() calls g()

- f is caller and g is callee

Steps of control flow

1. Setup parameters
2. Trf ctrl to callee
3. Setup local var
4. Store any results
5. Return ctrl to caller

Issues

Control Flow

- Need to jump to functional body when callee called
- Need to resume to next instruction in caller after done

Data storage

- Need to pass parameters to function
- Need to capture return result
- May have local variables

Additional

- May lead to overriding of data in caller by callee (interference)
- Calling g() multiple times may lead to insufficient space and overriding

Stack memory

Memory to store function invocation

**Stack Pointer** - Indicates the first free location in the stack region

**Frame Pointer** - Points to the frame and is used for traversing around the stack easily

On executing function call:

- **Caller:** Pass arguments with registers and/or stack ✓
- **Caller:** Save Return PC on stack
- **Transfer control from caller to callee**
- **Callee:** Save registers used by callee. Save old FP, SP
- **Callee:** Allocate space for local variables of callee on stack
- **Callee:** Adjust SP to point to new stack top

On returning from function call:

- **Callee:** Restore saved registers, FP, SP
- Transfer control from callee to caller using saved PC
- **Caller:** Continues execution in caller

Information needed for function invocation - Stack frame

- Return address of caller
- Arguments for the function
- Local variables
- Stack and frame pointer of caller
- GPR values (register spilling)

Callee stack frame will be on top of the caller

Dynamic memory (Heap)

Memory that the program/user specifies manually (eg. malloc, new)

Problems:

- Allocated only at runtime
  - Size not known at program compilation time
  - Cannot specify a region in data
- No definite deallocation timing
  - Must be freed explicitly by the program
  - Cannot place in stack region

Solution:

Add a region "Heap" for dynamic allocation

Problems with heap memory:

- Generation of holes in between data due to variable deallocation timing

OS context

Process identification

Features:

- Distinguish processes from each other (Unique)
- Communicated to the hardware

Process state

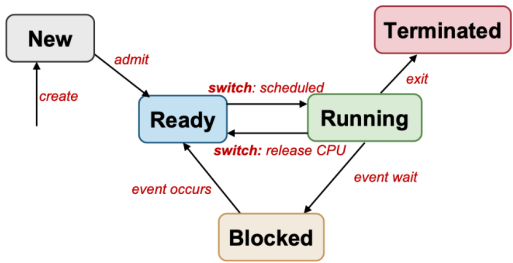
**New** New process that may still be under initialization (not schedulable)

Ready

**Running** Executed on CPU

**Blocked** Sleeping/waiting for event - cannot execute until event available

**Terminated** Finished execution and requires OS cleanup



Process control block

Table representing all processes containing entire execution context stored in OS memory

- PC, FP, SP and other GPR (only updated when swapped out)
- Memory region info (Text, data, heap and stack)
  - It is a "point" to real memory space **NOT actual memory space used by process**
- PID and process state

Context switching

1. Interrupt or syscalls
2. Save state of old process into PCB
3. Reload state of new process from PCB

- Involves a change to kernel mode and back

<p><b>Exceptions and interrupts</b></p> <p>Exceptions</p> <ul style="list-style-type: none"> <li>• Synchronous (due to program execution)</li> <li>• Machine level instructions arise errors</li> <li>• Exception handler executed automatically in software</li> </ul> <p>Interrupts</p> <ul style="list-style-type: none"> <li>• Asynchronous (Can happen anytime)</li> <li>• External events that cause execution to fail (hardware related errors)</li> <li>• Program execution suspended and <b>interrupt handler</b> executed automatically</li> </ul> <p><b>Instruction execution</b></p> <ol style="list-style-type: none"> <li>1. Read byte from PC and decode instruction</li> <li>2. Read 2 bytes to get the address/operands</li> <li>3. Perform ALU operations</li> <li>4. Store result into destination</li> <li>5. Check if any interruptions</li> </ol> <p>Interrupts can happen at anytime, and will remain pending until step 5 where it is handled</p> <p><b>Interruption handling</b></p> <ol style="list-style-type: none"> <li>1. Push PC and status register into hardware</li> <li>2. Disable interrupts</li> <li>3. Read <b>Interrupt Vector Table</b> - Table where the OS stores address of all interrupt handlers</li> <li>4. Switch to kernel mode</li> <li>5. Set PC to handler address and execute the instructions</li> </ol> <ul style="list-style-type: none"> <li>• OS populates the IVT table with address of interrupt routines</li> <li>• Hardware reads IVT to locate the handler</li> </ul>	
---	--

<p><b>System calls</b></p> <p><b>Application Program Interface</b> - Provides way of calling facilities/services in kernel</p> <p>Synchronous function only done in kernel mode providing system services</p> <p>Typically more expensive than library calls due to context switching</p> <p><b>Functions</b></p> <p><b>Process control</b> Direct processes like creating, terminating or synchronisation</p> <p><b>File/device manipulation</b></p> <p><b>Information maintenance</b> Get system data, time, etc</p> <p><b>Communication</b> Create, delete connections and send/receive messages</p> <p><b>Method (Programming language specific)</b></p> <ul style="list-style-type: none"> <li>• Library version with the same name and same arguments</li> <li>• User friendly library version</li> <li>• Using func <i>long syscall(long number);</i></li> </ul> <p><b>Mechanism</b></p> <ol style="list-style-type: none"> <li>1. User invoke library call</li> <li>2. Place call number in the designated location</li> <li>3. Library call executes a special instruction (<b>TRAP/syscall</b>) to change user to kernel mode</li> <li>4. (in kernel) syscall handler is determined (by a <b>dispatcher</b>)</li> <li>5. syscall handler is executed</li> <li>6. Syscall handler ends and control returned to the library call</li> <li>7. Return to user mode and continue normal function mechanism</li> </ol>	
--	--

<p><b>03. Process abstraction in Unix</b></p> <p>Process information</p> <p><b>Pid</b></p> <p><b>Proc state</b> Running, sleeping, stopped, zombie</p> <p><b>Parent pid</b></p> <p><b>Cumulative CPU time</b> For scheduling</p> <p><b>fork()</b></p> <p>Process creation</p> <p><b>Package</b> unistd.h and sys/types.h</p> <p><b>return</b> PID of newly created process(parent) and 0 (child process)</p> <p><b>Behaviour</b></p> <ul style="list-style-type: none"> <li>• Creates a child process</li> <li>– <b>Copy</b> data of parent (Independent memory space)</li> <li>– Unless explicitly modifying value at address, parent and child variables are distinct</li> <li>– Sane code, same address space</li> <li>– Differs by pid, ppid and fork() return value</li> </ul> <p><b>Implementation</b></p> <p>Clone the parent process</p> <ol style="list-style-type: none"> <li>1. Create address space of child process</li> <li>2. allocate new pid to child and pass to parent</li> <li>3. Create kernel process data structures</li> <li>4. Copy kernel environment of parent process</li> <li>5. Initialize child process context (pid, ppid, cpu_time = 0)</li> <li>6. Copy memory regions from parent <ul style="list-style-type: none"> <li>• Very expensive operation</li> <li>• Code, data, stack</li> </ul> </li> <li>7. Acquire shared resources</li> <li>8. Initialize hardware context for child process (copy parent registers)</li> </ol>	
---	--

<p>Problem: Malloc is very expensive operation</p> <p>Solution: Copy on write</p> <ul style="list-style-type: none"> <li>• Only duplicate a memory location when it is written to</li> </ul> <p><b>exec(@param)</b></p> <p>Replace current executing process image</p> <ul style="list-style-type: none"> <li>• Code and data replaced</li> <li>• PID/PPID intact</li> <li>• If fail, program just continues to run the next instruction</li> </ul> <p>Format</p> <p><b>param</b> char *path, char *arg0...</p> <ul style="list-style-type: none"> <li>• Note that last term MUST be <b>NULL</b> indicating end of argument list</li> </ul> <p><b>header</b> unistd.h</p> <p><b>exit(status)</b></p> <p><b>return</b> Does not return anything</p> <p><b>param</b> int status - returned to the wait call</p> <ul style="list-style-type: none"> <li>• Most system resource used by process are released on exit</li> <li>• return from main() implicitly calls exit(0)</li> <li>• Basic processes are not releasable <ul style="list-style-type: none"> <li>– Pid and status</li> <li>– Process accounting info</li> </ul> </li> </ul> <p><b>wait(&amp;status)</b></p> <p>Parent child synchronisation</p> <p><b>param</b> &amp;status - address to put return value</p> <p><b>header</b> sys/types.h and sys/wait.h</p> <p><b>return</b> pid of terminated process</p> <ul style="list-style-type: none"> <li>• Call is blocking - suspend operation until at least one child terminates</li> <li>• Cleans up remainder of child system resources (PID, status)</li> <li>• waitpid - used for waiting for specific child process</li> </ul>	
---	--

<p><b>Orphan and zombie process</b></p> <p><b>Zombie</b> - All processes that has exited but parent did not call wait</p> <p><b>Orphan</b> - Child process whose parent has been terminated</p> <ul style="list-style-type: none"><li>• Parenthood will be propagated up to init which may use wait() to clean up</li></ul> <p><b>04. Inter Process Communication Mechanism</b></p> <ul style="list-style-type: none"><li>• Shared memory</li><li>• Message passing<ul style="list-style-type: none"><li>– Pipes</li><li>– Signal</li></ul></li></ul> <p><b>Shared memory</b></p> <p>Communication through read/write to shared memory</p> <p>Advantages</p> <p><b>Efficient</b> Only require OS to setup shared region once</p> <p><b>Ease of use</b> Simple reads and write to memory</p> <ul style="list-style-type: none"><li>• Implicit communication</li><li>• Can store any type of information</li></ul> <p>Disadvantages</p> <p><b>Limited to single machines</b> Less efficient over different system</p> <p><b>Need Sync</b> Might have data races without synchronisation</p> <p><b>Race condition</b> - System behaviour is dependent on the context/interleaving of process → unpredictable outcome</p> <p><b>Steps</b></p> <ol style="list-style-type: none"><li>1. Create/locate a shared memory region M</li><li>2. Attach M to process memory space</li><li>3. Read from/write to M</li><li>4. Detach M from memory space after use</li><li>5. Destroy M</li></ol>	<ul style="list-style-type: none"><li>• Only one process need to do this</li><li>• Can only destroy if M is not attached to any process</li></ul> <p><b>Message Passing</b></p> <p>Explicit communication through exchange of message</p> <p><b>Naming</b> Have to identify the parties in the communication</p> <p><b>Synchronisation</b> Behaviour of sending/receiving operations</p> <ul style="list-style-type: none"><li>• Messages have to be stored in kernel memory space</li><li>• All sending/receiving operations have to be done through syscalls</li></ul> <p><b>Direct communication</b></p> <p>Sender/receiver explicitly name parties in communication</p> <ul style="list-style-type: none"><li>• One link per pair of communicating processes</li><li>• Need to know identity of other party</li></ul> <p><b>Indirect communication</b></p> <p>Message storage (<b>Mailbox/Port</b>)</p> <ul style="list-style-type: none"><li>• Can be shared among a number of processes</li></ul> <p><b>Synchronisation behaviour</b></p> <p>Blocking</p> <ul style="list-style-type: none"><li>• send and receive is blocked until a message has received/sent (other party is ready)</li></ul> <p>Non-blocking (asynchronous)</p> <ul style="list-style-type: none"><li>• execute immediately and sends information to somebody OR returns empty handed</li></ul> <p>Typically, receive is synchronous and send is asynchronous BUT send just buffers the message and only sends the message when the receiver is receiving (no loss)</p> <p>Message buffers</p> <ul style="list-style-type: none"><li>• Under OS control</li></ul>	<ul style="list-style-type: none"><li>• Decouples sender and receiver - less sensitive to variation in execution</li><li>• Mailbox capacity declaration in advance</li></ul> <p><b>Rendezvous</b></p> <ul style="list-style-type: none"><li>• Synchronous message passing</li><li>• Sender is blocked until receiver sends matching receive</li><li>• No buffering needed</li></ul> <p><b>Pros</b></p> <ul style="list-style-type: none"><li>• Applicable beyond a single machine</li><li>• <b>Portable</b> - Easily implemented on many platforms and processing environments</li><li>• Easier synchronisation<ul style="list-style-type: none"><li>– Implicit via send/receive behaviour</li><li>– Communication and synchronisation is done simultaneously</li></ul></li></ul> <p><b>Cons</b></p> <p><b>Inefficient</b> Usually requires OS intervention every operation</p> <p><b>Difficult to use</b> Requires information packing into specific format</p> <p><b>Unix Pipes</b></p> <p>3 different communication channels to user</p> <ol style="list-style-type: none"><li>1. stdin - standard in bounded to keyboard input</li><li>2. stderr - used for error messages</li><li>3. stdout - linked to screen</li></ol> <p>” ” symbol for linking input/output channels of each process to each other</p> <p>Create communication channel with 2 ends (one reading, one writing)</p> <ul style="list-style-type: none"><li>• Producer-consumer relationship</li><li>• Like anonymous file reading information FIFO</li></ul>	<p><b>Synchronisation</b></p> <ul style="list-style-type: none"><li>• Circular bounded byte buffer with implicit synchronization</li><li>• Writer wait when buffer is full</li><li>• Reader wait when buffer is empty</li><li>• Unidirectional (one write one read exclusively) vs bidirectional (any end for reading and writing)</li></ul> <p><b>pipe(@param)</b></p> <p><b>param</b> array of file descriptors</p> <p><b>return</b> 0 for success, != for errors</p> <p><b>Unix signals</b></p> <p>Asynchronous notification about an event sent to process/threads</p> <ul style="list-style-type: none"><li>• Handle signal via default or user supplied handlers</li><li>• Can only call async safe functions (signal/wait cannot be called in POSIX)</li><li>• eg. Kill, stop, continue, error</li></ul>
--	---	--	---

<h2>05.Process scheduling</h2> <p>Concurrency</p> <ul style="list-style-type: none"> <li>Multiple process progress at the same time</li> <li>Virtual parallelism (threading, context switching)</li> <li>Physical parallelism (multicore, multi CPU)</li> </ul> <h3>Scheduling</h3> <p>More ready processes vs limited CPU</p> <ul style="list-style-type: none"> <li>Each process diff CPU usage</li> <li><b>Scheduling algorithms</b> - Deciding the process to run based on environment and process behaviour</li> <li>Processes have phases of io-bound and cpu intensive activities</li> </ul> <p>Processing environment (sorted by increasing user interaction and response times):</p> <ol style="list-style-type: none"> <li>Batch</li> <li>Interactive</li> <li>Real time</li> </ol> <h3>Objectives</h3> <p><b>Fairness</b> Should get fair share of CPU time n prevent <b>starvation</b> - processes never have access to CPU</p> <p><b>Utilization</b> Hardware is used all the time</p> <p><b>non-preemptive</b> - Process gives up resources voluntarily <b>pre-emptive</b> - Process given a fixed quota to run and is suspended for other processes</p> <h3>Timeline</h3> <ol style="list-style-type: none"> <li>Scheduler is triggered (OS take over)</li> <li>Context switching can happen and current process information stored and placed back in queue</li> <li>Pick process to run using algorithm and setup its context</li> </ol> <h3>Batch</h3> <p>Non-preemptive used dominantly cos minimal user interaction needed</p>	<h3>Objectives</h3> <p><b>Turnaround time</b> Total time till finish running <b>including waiting</b> (finish time - arrival time)</p> <p><b>Throughput</b> Rate of task completion</p> <p><b>Makespan</b> Total time to complete <b>ALL</b> tasks</p> <p><b>CPU utilization</b> Percentage of <b>time</b> CPU is used</p> <h3>Algorithms</h3> <ol style="list-style-type: none"> <li>First Come First Serve (FCFS) <ul style="list-style-type: none"> <li>Scheduled using FIFO queue based on arrival time (bad turnaround time)</li> <li>Guarantee no starvation - number of tasks in front of task is strictly decreasing</li> <li>Reordering can reduce waiting time</li> <li><b>Convoy Effect</b> - Not all hardware resources are used concurrently because a process blocks it</li> </ul> </li> <li>Shortest Job First (SJF) <ul style="list-style-type: none"> <li>Choose tasks with smallest total CPU time</li> <li>Estimate CPU time for tasks in advance using the previous CPU-bound phases</li> <li><math>Estimated_{n+1} = \alpha Actual_n + (1 - \alpha) Predicted_n</math></li> <li>Starvation is possible as long tasks end up continuously waiting</li> </ul> </li> <li>Shortest Remaining Time (SRT)</li> </ol> <h3>Interactive system</h3> <h3>Objectives</h3> <p><b>Response time</b> Time between request and response by system</p> <p><b>Predictability</b> Decreased variation in response time</p> <h3>Preemptive scheduling</h3> <ul style="list-style-type: none"> <li>Periodic scheduler with time interrupts</li> <li><b>Time quantum</b> - Execution duration given to a process <ul style="list-style-type: none"> <li>Constant(regardless whether the process has given up processing halfway) OR Variable(full time quantum given to a process)</li> <li>-</li> </ul> </li> </ul>	<h3>Algorithms</h3> <ol style="list-style-type: none"> <li>Round Robin (RR) <ul style="list-style-type: none"> <li>Store tasks in FIFO queue and pick tasks to run until <ul style="list-style-type: none"> <li>(a) Fixed time slice elapsed</li> <li>(b) Task gives up CPU voluntarily</li> <li>(c) Task blocks</li> </ul> </li> <li>Task placed at end of queue to wait for next turn</li> <li>Response time guaranteed bounded by num task time quantum</li> <li>Timer interrupt needed</li> <li>Choice of time quantum duration is important <ul style="list-style-type: none"> <li>Longer quantum: Better utilization but greater wait time</li> <li>Shorter quantum: Bigger overhead/ lower utilization but shorter wait time</li> </ul> </li> </ul> </li> <li>Priority scheduling <ul style="list-style-type: none"> <li>Assign tasks processes priorities based on impt and select them</li> <li>Starvation for low priority tasks <ul style="list-style-type: none"> <li>Solution: <ul style="list-style-type: none"> <li>Lower priority every time the process have executed</li> <li>Give processes a time quantum</li> </ul> </li> </ul> </li> <li><b>Priority Inversion</b> - Lower priority tasks locks the resources used by a higher priority task leading to deadlock OR allowing lower priority tasks to run before <ul style="list-style-type: none"> <li>Solution: <ul style="list-style-type: none"> <li>Temporarily increase priority of tasks that lock resources until it unlocks</li> <li>Low priority task inherit the priority of high priority tasks which is restored once unlocked</li> </ul> </li> </ul> </li> </ul> </li> <li>Multi-level Feedback Queue (MLFQ) <ul style="list-style-type: none"> <li>Runs processes with higher priorities but lower priority when the job fully utilise its time quantum</li> <li>Processes that voluntarily gives up/blocks before time quantum retains its priority</li> <li>New processes have highest priority</li> <li>Processes with same priority run in RR</li> <li>Long CPU processes are starved as its priority becomes very low and cannot complete <ul style="list-style-type: none"> <li>Solution: Occasionally reset tasks to full priority</li> </ul> </li> </ul> </li> </ol>
--	--	--



- Processes can exploit and voluntarily give up to retain its priority
  - Solution: Peg priority based on total CPU usage instead

4. Lottery scheduling

- Give out tickets to processes for various system resources and randomly pick winners to be granted the resources
- Responsive** - Newly created processes can immediately join the lottery
- Good control - Process can be assigned quantity of tickets and resources which can be shared among its children/threads
- Simple implementation

06.Synchronisation primitives

Race condition

- Execution of sequential processes should be **deterministic** - Repeated execution returns the same result... however,
- Process sharing modifiable resources AND execute concurrently by interleaving may cause synchronization problems - non deterministic outcomes
- Order determines the execution outcome

Critical section

Region where unsynchronised access could lead to incorrectly interleaving scenarios  
Only **ONE** process should be in the critical section at one time

Properties

**Mutual exclusion/Mutex** All other processes should be blocked from entering critical section if occupied

**Progress** If no process occupying section, waiting process should be granted access

**Bounded wait** Upper bound on number of times other processes can enter critical section before a waiting process

**Independence** Process not in critical section should not block other process

Incorrect synchronization

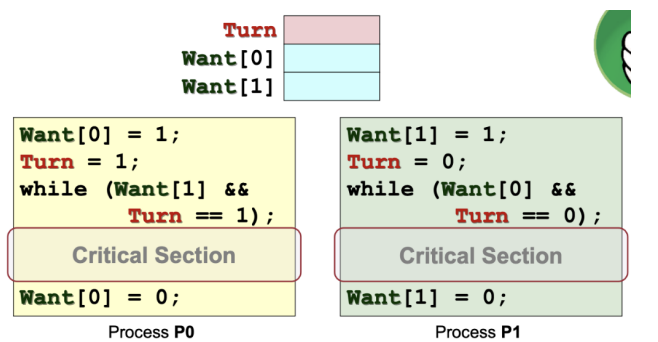
**Deadlock** All processes blocked

**Livelock** Processes keep changing states to avoid deadlock resulting in no progress (Deadlock avoidance mechanism)

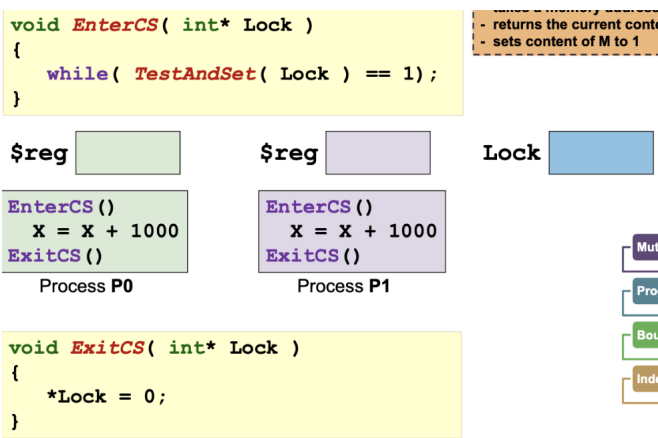
**Starvation**

Implementations

- High level language implementation (Peterson Algorithm)
  - Busy waiting - Repeatedly tests while-loop condition instead of going into blocked state (Wastes CPU power)
  - Low level - Error prone and complicated
  - Not general - cannot extend to allowing multiple access instead of Mutex



- Assembly implementation
  - Atomic instruction to load value and replace value in memory location (**TestAndSet**)
  - Atomicity** - Prevents race condition of process viewing and upon modification, the value changes
  - No bounded wait unless fair scheduling algorithm employed
  - Employs busy waiting still



- High level synchronization mechanism (Semaphore)
  - Generalised mechanism to block a number of processes (sleeping),  $S$ , and unblock sleeping processes
  - wait()** - Blocks if  $S \leq 0$ , decrement  $S$

- signal()** - Increment  $S$ , wake up sleeping processes, operation **never** blocks
- Invariant** -  $S_{current} = S_{initial} + N_{signal} - N_{wait}$

Conditional variables

- Tasks wait for certain events to happen
- Event completion/begin is broadcasted

POSIX semaphore

header semaphore.h

wait pthread\_mutex\_unlock

signal pthread\_mutex\_lock

Midterm

- Tail-call optimization** - Stack frame can be replaced in iterative recursion since all information is retained in the 'return' function call at every iteration

07.Threads

Motivation

- Processes are expensive
  - Process creation via fork() duplicates memory space and process context
  - Requires context switching and saving/restoring process information repeatedly
- Hard to communicate with each other
  - Independent memory space (IPC communication only)
- Need easy way to execute some instructions simultaneously (more threads of control to same process)

Features

- Multithreaded process** - Single process with multiple threads
- Threads share the same **Memory context**(Text, data, heap) and **OS context**(PID)
- Contains unique information
  - Identification (thread id)
  - Registers (GPR and special)
  - Stack

**Benefits**

**Economy** Less resources to manage

**Resource sharing**

**Responsiveness**

**Scalability** Can take advantage of multiple CPUs

**Problems**

- Syscall concurrency
  - Parallel execution of threads may result in parallel syscalls
- Process behaviour
  - Unable to impact process operations

**Thread implementation**

1. User thread
  - Implemented as a user library - More flexible and configurable
  - Runs on any OS
  - Kernel unaware of threads in process (scheduling at process level)
  - If process is blocked, all the threads in the process cannot run
2. Kernel thread
  - Implemented in the OS through syscalls (slower and more resource intensive)
  - Scheduling among threads instead of via process
  - Use threads in kernel operations
  - Less flexible
    - Used by all multithreaded programs so kernel must cater to all
    - Balance btw many and less features
3. Hybrid thread
  - OS schedules on kernel threads
  - User threads can be bound to kernel threads
  - If the user thread is blocked, the kernel thread that it is assigned to is blocked and the OS will execute another kernel thread
  - Lead to greater flexibility

**Posix Threads (pthreads)**

**pthread\_create**

**param** tid, attributes, &function, function\_params

**return** 0 success, != 0 fail

**pthread\_exit**

Automatically called after function finishes

**param** exit\_value

**behaviour** return value of function will be the "exit value"

**pthread\_join**

**param** tid, **&status** - Exit value returned by target pthread

**08.Synchronization classics**