

01. Stable matching

Both sides rank each other and goal is to pair each up **Constraints** No **rogue couples** - matched partner likes someone more than current partner that also likes them back

Gale-Shapley Algo

```
Initialize each person to be free.
while (some man is free and hasn't proposed to every woman) {
    Choose such a man m
    w = 1st woman on m's list to whom m has not yet proposed
    if (w is free)
        assign m and w to be engaged
    else if (w prefers m to her fiancé m')
        assign m and w to be engaged, and m' to be free
    else
        w rejects m
}
```

Invariant

- If woman not on boy list, she has a better current fav
- Boy choice is strictly worsening
- Girl choice cannot worsen (weakly increasing)

02. Algorithm Analysis

Algorithm Finite sequence of well-defined instruction to solve computational problem

- Optimize running time
- Runtime is machine and input dependent

Word-RAM model

- Word is basic unit of memory (few bytes)
- Runtime \approx number of instructions taken
- Ram can be accessed in the same time irrespective of location
- Every operation involves constant number of words and cycles by CPU

Problem vs Algorithm

- Problem does not specify approach but algorithm defines ways/methods to solve prob

Polynomial time

- There exist constants $c, d \in \mathbb{R}$ on every input of size N such that its running time is bounded by cN^d
- Brute Force - Checks all possible solution and typically takes exponential time 2^N
- Generally efficient because c and d are low

Worst-Case Analysis

- Bound on largest possible running of algorithm on input
- Captures efficiency in practice

vs Average case running time

- Hard to accurately model real instances by random distributions
- Tuned for certain distributions

Asymptotic Order of Growth

- Compare for asymptotically large values of input sizes (rate of growth)
- Suppress constant factor and lower-order terms

Bounds

Upper Bound $T(n) = O(f(n))$: Exist constants $c, n_0 \geq 0$ such that $\forall n \geq n_0$, we have $T(n) \leq cf(n)$

Lower Bound $T(n) = \omega(f(n))$: Exist constants $c, n_0 \geq 0$ such that $\forall n \geq n_0$, we have $T(n) \geq cf(n)$

Tight Bound $T(n) = \theta(f(n))$: $T(n)$ is both $O(f(n))$ and $\omega(f(n))$

Properties

Transitivity If $f = O(g)$ and $g = O(h)$ then $f = O(h)$

Additivity If $f = O(h)$ and $g = O(h)$ then $f + g = O(h)$

Bounds for some functions

- Polynomials - Just take the highest power
- Logarithms - Base doesn't matter and always smaller than polynomial
- Exponential - Always grows faster than polynomial

Common runtimes

- $O(n)$ - List traversal or merging lists
- $n \log n$ - sorting, intervals, divide and conquer algos (recursive)
- n^2 - Finding pairs of elements
- n^k - Enumerate subsets of k nodes, set disjoint
- Exponential - All subsets

03. Graph

- Describes relationships between nodes/entities

Representation

Adjacency Matrix $N \times N$ matrix where $A_{uv} = \text{weight}_{u,v}$ if (u,v) is an edge

- Good for big graphs/a lot of edges(dense)
- $O(1)$ time for checking edge
- Space is $O(n^2)$
- Identifying all edges take $O(n^2)$ time

Adjacency list List of linked lists representing nodes that are attached to each node

- Good for sparse graphs/trees
- $O(\deg(u))$ for checking edge
- Space is $V+E$
- Identifying all edges takes $O(V+E)$ time

Path and Connectivity

- **Path** - Sequence $P \in \{v_0, v_1, \dots, v_k\}$ of nodes such that each adjacent pair is joined by an edge

Simple All nodes are distinct

Cycle Path such that $v_0 = v_k, k > 2$ and first k-1 nodes are distinct

- **Connected** - There is a path for all pairs of nodes
- **Tree** - Undirected connected graph that does not contain a cycle
 - Any 2 implies Tree
 - G is connected
 - G does not contain cycle
 - G has n-1 edges
- **Rooted tree** - Arrangement of tree such that root is top and nodes attached to parent is in different level

Connectivity problems

s-t connectivity Is there path between any 2 points s, t

s-t shortest path prob Length of shortest path btw s, t

Breadth First Search Algorithm

- Explore outward from some node, s in all directions adding nodes one layer at a time
- Algorithm
 1. $\text{Layer}_0 = s$
 2. $L_1 = \text{neighbours of } L_0$
 3. $L_n = \text{neighbours of } L_{n-1}$ not in all prev layers

Properties

- Induction theorem: $\forall i, L_i$ consist of nodes at $\text{dist}(s) = i$
- Let T be a BFS tree of $G = (V, E)$, and let (x, y) be an edge of G. Then the level of x and y differ by at most 1.
- BFS runs in $O(V + \sum E)$ given adjacency list

Bipartite graph

- Independent sets such that no nodes within the set are connected to each other but are connected to nodes of other sets
- Colour each node red or blue and every edge has strictly one red and blue end each
- Bipartite graphs cannot contain odd length cycles

Algo for Bipartiteness

- Let G be a connected graph and $L_0..L_k$ are layers produced by BFS
 1. Edge between 2 nodes in same layer
 - Graph is bipartite as we can colour alternate layers as red and blue
 2. No edge btw nodes in same layer
 - Graph cannot be bipartite as there is odd length cycle
 - Let x,y be the 2 points in same level and has edge btw them
 - Cycle: s-x-y-s (s-x == y-s) Odd length

Directed graph

- Asymmetric edges: (u,v) (edge from u to v) does not imply (v,u)
- Directed reachability problem: Find all nodes reachable from node s

Strong Connectivity

Mutually reachable Path from u to v and vice versa

Strongly connected Every pair of nodes are mutually reachable

- Strongly connected iff every node reachable from s and s reachable from every node

Algorithm Pick any node and run bfs on G and G^{rev} (reverse orientation of every edge in G)

- G is strongly connected if all nodes reached in both BFS executions
- Runs in $O(m+n)$ time

Directed Acyclic Graph

- Directed graph that has no directed cycle
- **Precedence constraints** - Edge (v_i, v_j) means v_i must precede v_j
- **Topological order** - For every edge (v_i, v_j) we have $i < j$
 - If G has a topological order, then G is a DAG (Proof by contra)
 - If G is a DAG then G has a node without incoming edges (Proof contra)

04. Greedy algorithms

- Consider jobs in some order and immediately take the job that is compatible with the previous jobs
- Has some heuristic to sort the jobs

Scheduling intervals

- Maximize number of jobs run
- Heuristic: end time, start time, interval size, fewest conflict

Interval Partitioning

- Minimize number of classrooms to run lectures concurrently
- Heuristic: start time
- Assign lecture to first compatible classroom

Minimize lateness scheduling

- Heuristic: Earliest deadline, Shortest processing time first, Smallest slack ($d_j - t_j$)
- Start with an optimal solution and inductively reduce number of inversions until it reaches the greedy solution

Inversions

- Pair of jobs i and j where $d(i) > d(j)$ but j schedule before i
- Greedy solution does not have any inversions

Greedy analysis Strategies

1. Each step in greedy solution is at least as good as other solutions
2. Exchange argument: Transforming solution to greedy algorithm without hurting its quality
3. Structural: Every solution has a certain value and greedy algos can be in this bound

Optimal offline caching

- Eviction strategy that minimises number of cache misses
- Note that the full sequence of requests is known a priori
 - vs Online: Request not known in advance
- Heuristic: Farthest in future
 - Evict item in cache that is not requested until farthest in the future

05. Divide and conquer algo

- Break up problem into several parts and sort each part recursively
- Combine solutions to sub-problems into overall solution

$$T(n) \leq T(\lceil \frac{n}{2} \rceil) + T(\lfloor \frac{n}{2} \rfloor) + n$$

07. Network Flow

Minimum Cut/Max Flow Problem

Environment

- Directed graph with "capacity" as edge weights $c(e)$
- 2 nodes denoted as source and sink

s-t cut

- Partition (A,B) of graph V with $s \in A$ and $t \in B$
- Capacity of a cut: $cap(A,B) = \sum_{edges\ out\ of\ A} c(e)$

s-t flow

- Function that satisfies:
 - $\forall e \in E 0 \leq f(e) \leq c(e)$
 - $\forall v \in V - (s,t), \sum_{e\ into\ v} f(e) = \sum_{e\ out\ of\ v} f(e)$

- Value of flow: $v(f) = \sum_{e\ out\ of\ s} f(e)$

Goals

Min s-t cut Find an s-t cut of minimum capacity

Max s-t flow Find an s-t flow of maximum value

Flow Value Lemma

- Let f be any flow, and let (A,B) be any s-t cut. The net flow sent across the cut is equal to the amount leaving s

$$\sum_{e\ out\ of\ A} f(e) - \sum_{e\ into\ A} f(e) = v(f)$$

Weak Duality


- Let f be any flow. Then for any s-t cut (A,B) we have $v(f) \leq cap(A,B)$

Certificate of Optimality

- Let f be any flow and let (A,B) be any cut, if $v(f) = cap(A,B)$ then f is a max flow and (A,B) is a min cut

Residual Graph


Original edge: $e = (u, v) \in E$.
• Flow $f(e)$, capacity $c(e)$.



Residual edge.

- "Undo" flow sent.
- $e = (u, v)$ and $e^R = (v, u)$.
- Residual capacity:

$$c_f(e) = \begin{cases} c(e) - f(e) & \text{if } e \in E \\ f(e) & \text{if } e^R \in E \end{cases}$$



- Residual graph is all residual edges with positive residual capacity
- $E_f = \{e : f(e) < c(e)\} \cup \{e^R : c(e) > 0\}$

Ford Fulkerson

1. Find an augmenting path/flow
 - Can use BFS or DFS to find the path and get the bottleneck edge
 - Flow value is the smallest value
 - Keep track of the flow value
2. Augment the graph using residual edges
3. Repeat until no feasible paths possible

Proof

- Flow f is a max flow iff there are no augmenting path
 - (i) There exists a cut (A,B) st $v(f) = cap(A,B)$
 - (ii) Flow f is a max flow
 - (iii) There is no augmenting path relative to f
- (i)-(ii): Corollary to weak duality to weak duality lemma
- (ii)-(iii): Contrapositive
 - Suppose there exists an augmenting path, then we can improve flow by sending flow along path
- (iii)-(i)
 1. Let f be a flow with no augmenting path
 2. Let A be set of vertices reachable from s in residual path
 3. By definition of A, $s \in A$
 4. By definition of f, $t \notin A$

Runtime

- All capacities are integers btw 1 and C
- Every flow value $f(e)$ and every residual capacity remains an integer throughout the algorithm
- Algorithm terminates in at most $v(f^*) \leq nC$ iterations
 - Each augmentation increase value by at least 1
- If C = 1, Ford-Fulkerson runs in O(mn) time
- If all capacities are integers, then there exists a max flow f for which every flow value $f(e)$ is an integer.
 - Since algorithm terminates, theorem follows from point 2